

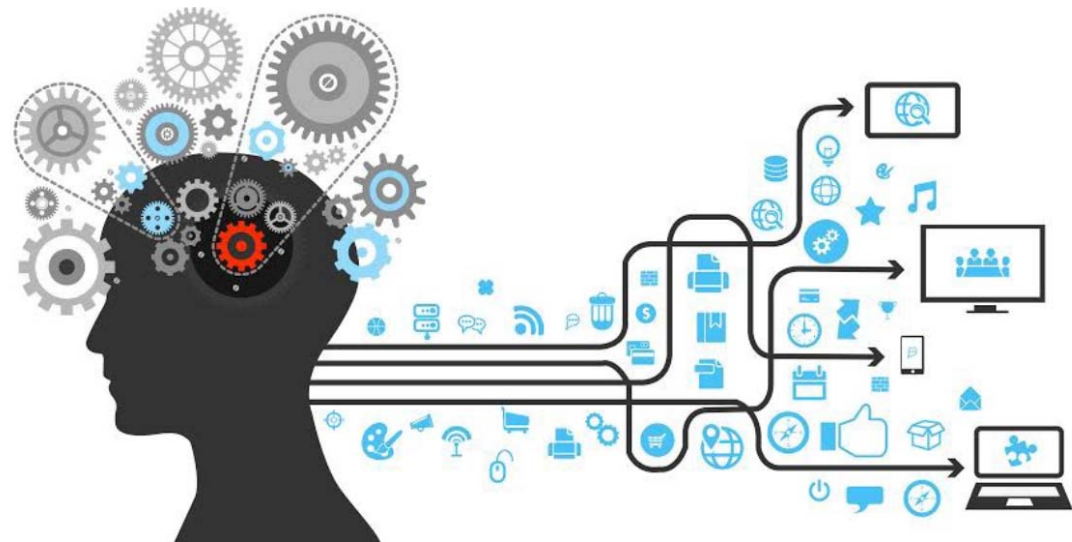
Computer Vision


Early vision: Just one image

School of Electronic & Electrical Engineering


Sungkyunkwan University

Hyunjin Park





Linear Filters



Representing image regions

- We can generate a number of intrinsic images from a given image
 - Depth/disparity
 - Surface albedo/color
 - Surface normal
 - ...

- How can we organize these into surfaces?
 - Identify attributes of regions
 - Bounding edges
 - Texture
 - Spatial aggregations of pixels
 - Segmentation



Taxonomy

- Signal processing
 - Discrete-time signal processing
 - Wavelet tour of signal processing
- Image processing
 - Two-dimensional signal and image processing
 - The Fourier transform and its applications
- Tools which have become indispensable to computer vision
 - Linear filters
 - Over-complete (pyramid) representation

Systems and filters

- Filtering

- *Form a new image* whose pixels are a combination of original pixel values

- Goals


- Extract useful information from the images
 - Features (edges, corners, blobs, ...)
- Modify or enhance image properties
 - Super-resolution, in-painting, de-noising

Systems and filters

De-noising



Super-resolution



In-painting



Image filtering

- Filtering

- Modify the pixels in an image
- Based on some function of a local neighborhood of the pixels

10	5	3
4	5	1
1	1	7

Some function



	7	

Linear filtering

- Linear case is the simplest and most useful
 - Form a new image by replacing each pixel with a weighted sum (i.e., linear combination) of its neighbors, using the same set of weights at each point
- The prescription for the linear combination is called the *kernel*

10	5	3
4	5	1
1	1	7

 $*$

0	0	0
0	0.5	0
0	1.0	0.5

 $=$

	7	

kernel

Moving average

- Example: 2D discrete-space moving average with 3×3 window

$$\begin{aligned} g[n, m] &= \frac{1}{9} \sum_{k=n-1}^{n+1} \sum_{l=m-1}^{m+1} f[k, l] \\ &= \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n-k, m-l] \end{aligned}$$

h

1	1	1
1	1	1
1	1	1

$\frac{1}{9}$

$$(f * h)[m, n] = \frac{1}{9} \sum_{k, l} f[k, l] h[m-k, n-l]$$

Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0								

Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10							

Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20						

Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20	30					

Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20	30	30				

Moving average

$F[x, y]$


0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

Moving average

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect
 - Remove sharp features



Linear filter properties

- Linear
 - Output is a linear function of the input
 - Superposition: $h * (f_1 + f_2) = (h * f_1) + (h * f_2)$
 - Scaling: $h * (kf) = k(h * f)$
 - $S[\alpha f_1 + \beta f_2] = \alpha S[f_1] + \beta S[f_2]$

Linear filter properties

- Shift-invariant

- Output is a shift-invariant function of the input

- Shift the input image two pixels to the left,
the output is shifted two pixels to the left

- If $f[n, m] \xrightarrow{S} g[n, m]$, then $f[n - n_0, m - m_0] \xrightarrow{S} g[n - n_0, m - m_0]$

Correlation filtering

- Size of the averaging window: $(2k + 1) \times (2k + 1)$


$$g[i, j] = \frac{1}{(2k + 1)^2} \sum_{u=-k}^k \sum_{v=-k}^k f[i + u, j + v]$$

- Generalize to allow different weights depending on neighboring pixel's relative position

$$g[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k h[u, v] f[i + u, j + v]$$

$$g = h \otimes f$$


Correlation filtering



Convolution filtering

- Convolution: $g = h \star f$
 - Flip the filter in both dimensions (bottom to top, right to left)
 - Then apply cross-correlation

$$g[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k h[u, v] f[i - u, j - v]$$



Convolution vs. correlation

- Convolution:

$$g[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k h[u, v] f[i - u, j - v]$$

$$g = h \star f$$

- Cross-correlation:

$$g[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k h[u, v] f[i + u, j + v]$$

Convolution vs. correlation

- Convolution:

- Integral that expresses the amount of overlap of one function as it is shifted over another function
- Convolution is a filtering operation

- Cross-correlation:

- Computes a measure of similarity of two input signals as they are shifted by one another
- The correlation result reaches a maximum at the time when the two signals match best
- Correlation is a measure of relatedness of two signals

Linear filter examples



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Linear filter examples



Original

0	0	0
0	0	1
0	0	0



Shifted left
by 1 pixel with
correlation


Linear filter examples



Original

 $\frac{1}{9}$

1	1	1
1	1	1
1	1	1



Blur (with a box filter)

Linear filter examples



Original

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$


1	1	1
1	1	1
1	1	1




Sharpening filter

Linear filter examples

Sharpening filter




before



after

Smoothing: Average filter

- We can *reduce noise by smoothing*
- Smoothing by averaging




Smoothing: Gaussian filter

- Average smoothing is not appropriate for a defocused lens
 - A single point of light viewed in a defocused lens looks like a fuzzy blob
 - Averaging process would give a little square
- Gaussian kernel gives a good model of a fuzzy blob



$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$

1	2	1
2	4	2
1	2	1




Smoothing: Gaussian filter

- Smoothing with a Gaussian




Smoothing: Gaussian filter

- Size of kernel




$\sigma = 5$ with 10×10 kernel $\sigma = 5$ with 30×30 kernel




Smoothing: Gaussian filter

- Variance of Gaussian
 - Determines extent of smoothing



$\sigma = 2$ with 30×30 kernel



$\sigma = 5$ with 30×30 kernel


Smoothing: Gaussian filter

■ Rows

- Smoothing with Gaussians of different width

■ Columns

- Different realizations of an image of Gaussian noise




Mean filter

- Smoothing by averaging
 - Good when the average is taken over a homogeneous neighborhood with zero-mean noise
 - When the neighborhood straddles the boundary between two homogeneous regions, the estimate results in blurring of the boundary


Median filter

- Median filter
 - Non-linear filter
 - Method:
 - 1) Rank-order neighborhood intensities
 - 2) Take middle value
 - No new gray level emerge




Median filter

- Median filter
 - Remove spikes
 - Good for impulse, salt & pepper noise
 - Less sensitive to outliers compared to mean filter



Salt and pepper noise Median filtered



Plots of a row of the image

Median filter

- Median filter
 - Not always optimistic

3 x 3 median filter



Sharpens edges, destroys edge cusps and protrusions

Median filter

- Median filter
 - Not always optimistic

Comparison with Gaussian filter



E.g.) Upper lip smoother, eye better preserved

Median filter

- Median filter
 - Not always optimistic



10 times 3 x 3 median filter



Patchy effect: Important details lost (e.g., earring)

Filters as templates

- Filtering the image
 - Applying a filter at some point can be seen as taking a dot-product between the image and some vector
 - Filtering the image is a set of dot products
- Filters look like the effects they are intended to find matched filters




Filters as templates


- Filters as templates
 - Can be used for template matching
 - Use normalized cross-correlation to find a given pattern in the image

$$\frac{\sum_{x,y}(A_{x,y} - A_{mean})(B_{x,y} - B_{mean})}{\sqrt{\sum_{x,y}(A_{x,y} - A_{mean})^2} \sqrt{\sum_{x,y}(B_{x,y} - B_{mean})^2}}$$

Template matching




Scene




Template

Template matching



Detected template



Correlation map

Template matching




Scene



Template

Template matching



Correlation map



Template